

Ultimate Source Of Truth

1. Logic Flow (Ingestion Service, No Storage Phase)

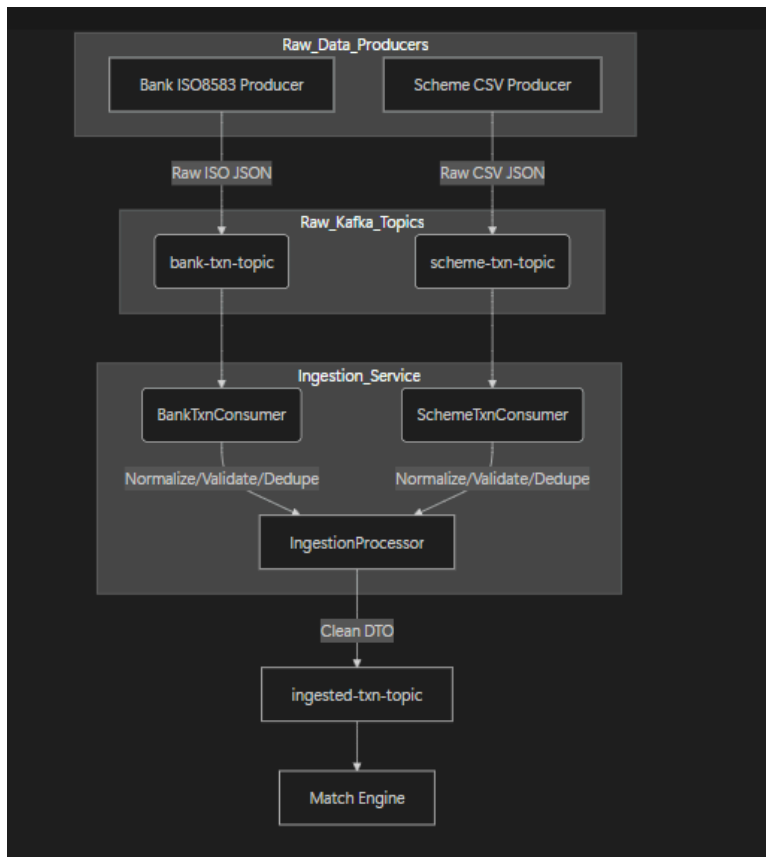
a. Input:

- **Source 1:** Bank/switch event (raw ISO8583, JSON, or other)
- **Source 2:** Scheme/clearing file line (raw CSV, JSON, etc.)

b. Process:

1. **Kafka Consumer receives raw message** from the appropriate topic.
 2. **Normalize raw message** → TxnRecordDTO
 - Parse, map, validate all required fields.
 - Set source-specific fields (`sourceType` , `batchId` , `schemeName` , etc.).
 3. **Validation & error handling:**
 - If invalid, log, update error metric, skip.
 4. **Deduplication:**
 - Filter out duplicate transactions by txnId.
 5. **Metrics update:**
 - Increment counters for total, deduped, errors, by source.
 6. **Emit normalized DTO:**
 - Pass validated, deduped TxnRecordDTO to the match-engine (next microservice).
 - This is done by **pushing to a Kafka topic** consumed by the match-engine.
- c. **No storage is performed here**—ingestion is stateless except for dedup memory/cache.
-

2. Mermaid Sequence Diagram



3. Service Handoff

- **What is passed on:**
 - **Type:** TxnRecordDTO
 - **Format:** As a JSON-serialized DTO, pushed to a Kafka topic (e.g., ingested-txn-topic or match-engine-in).
 - **Contents:**
 - All normalized fields (txnId , amount , currency , etc.)
 - sourceType (BANK_SWITCH or SCHEME_FILE)
 - Source-specific fields filled as needed (batchId , schemeName , etc.)
- **To whom:**
 - **Downstream consumer = Match Engine microservice**
 - Its job: fetch from this topic, perform matching, reconciliation, and *then* persist or further process.
- **What to know:**
 - **Ingestion service remains stateless.**
 - **No database/storage at this phase.**
 - **Metrics, dedup, and error handling all remain inside ingestion for observability only.**

This doc is the final truth for your ingestion-service phase-1.

Every engineer, PM, or reviewer can know the exact flow, what gets emitted, and what stays inside the service.

If you want a YAML/Markdown for doc, or a visual, say so—ready for copy-paste to a README or design doc.

▼ SOURCES

Field Name	Common?	BANK_SWITCH	SCHEME_FILE	Notes
txnId	✓ Common	✓ Required	✓ Required	Always present
cardNumber	✓ Common	✓ Required	✓ Required	Always present
amount	✓ Common	✓ Required	✓ Required	Always present
txnTimestamp	✓ Common	✓ Required	✓ Required	Always present
currency	✓ Common	✓ Required	✓ Required	Always present
merchantId	✓ Common	✓ Optional	✓ Optional	Usually present
terminalId	✓ Common	✓ Optional	✓ Optional	Usually present
responseCode	✓ Common	✓ Optional	✓ Optional	Usually present
channel	✗ Source-specific	✓ Often Present	✗ Not present/missing	e.g., "POS", "ATM"—bank only
batchId	✗ Source-specific	✗ Not present/missing	✓ Often Present	Scheme only (e.g., Visa batch/file ID)
schemeName	✗ Source-specific	✗ Not present/missing	✓ Often Present	Scheme only (e.g., "VISA")
authCode	✗ Source-specific	✓ Often Present	✗ Not present/missing	e.g., for authorizations
sourceType	✓ Common	"BANK_SWITCH"	"SCHEME_FILE"	Enum, always present
rawSourceRecord	✓ Common	✓ Optional	✓ Optional	Raw/original message

▼ sending

```

[Producer Java Object] DTO
|
| (Jackson/Gson: Object → JSON String)
v
[JSON String]
|
| (StringSerializer: String → Bytes)
v
[Kafka Topic (Bytes = JSON String)]
|
| (StringDeserializer: Bytes → String)
v
[JSON String]
|
| (Jackson/Gson: JSON String → Object)

```

Recon Ingestion Service: Source of Truth

1. What This Service Does

- **Acts as the “cleaning gateway”** for all transactions.
- **Consumes raw, dirty, and inconsistent transaction records** from upstream (bank switch/ISO8583 and scheme CSV).
- **Normalizes** them into a single DTO format.
- **Validates** every record (schema, constraints, required fields, data types).
- **Deduplicates** (using composite key: txnId + sourceType).
- **Feeds only clean, valid, deduped data** to the match engine via the “ingested-txn-topic.”
- **Tracks metrics** for all stages (total, success, deduped, error).
- **Separation of raw and clean data** (raw topics for forensics/auditing, clean topic for matching engine)

3. Components & Contracts

A. Producers

- *TestTxnProducer* creates intentionally dirty and clean bank/scheme messages.
- Pushes to `bank-txn-topic` and `scheme-txn-topic`.

B. Kafka Topics

- **bank-txn-topic**: Raw ISO8583-like, JSONified by producer.
- **scheme-txn-topic**: Raw CSV, JSONified by producer.
- **ingested-txn-topic**: Clean, validated, normalized, deduped DTOs.

C. Consumers

- **BankTxnConsumer**: Reads from `bank-txn-topic`, deserializes, hands to processor.
- **SchemeTxnConsumer**: Reads from `scheme-txn-topic`, deserializes, hands to processor.

D. IngestionProcessor

- **Normalization**: Maps raw fields to internal DTO (TxnRecordDTO).
- **Validation**: Uses Jakarta Validation to enforce constraints (not null, >0, proper types).
- **Deduplication**: Only allows unique (txnId + sourceType) to pass through.
- **Metrics**: Exposes counters for monitoring.

- **Downstream Producer:** Pushes to `ingested-txn-topic` for match engine.
-

4. Key Lessons / Gotchas Fixed

- **Deduping only on txnid is WRONG.** Must dedupe on *both* txnid and sourceType (bank vs scheme).
 - **Raw records are for forensics and reprocessing, not for main engine.**
 - **Validation failures and deduped records are never sent downstream.**
 - **All logic is service-injected, ObjectMapper as Bean, no static hackery in prod pipeline.**
 - **Metrics and logging added at all key points.**
 - **Easy to extend for distributed cache or persistent dedupe later.**
-

5. Reference Implementation Notes

- *BankTxnConsumer* & *SchemeTxnConsumer*:
Constructor-inject `ObjectMapper`, hand off to `IngestionProcessor`.
 - *IngestionProcessor*:
 - `ConcurrentHashMap.newKeySet()` for deduplication, composite key is `"txnid|sourceType"`.
 - All validation/dedup logic is centralized and logged.
 - **Pushes to `ingested-txn-topic` using the composite key as Kafka key.**
 - *TestTxnProducer*:
Standalone utility for simulating real traffic.
-